



GraphSAINT: Graph Sampling Based Inductive Learning Method

Hanqing Zeng^{*,1}, Hongkuan Zhou^{*,1}, Ajitesh
Srivastava¹, Rajgopal Kannan², Viktor Prasanna¹

1. University of Southern California

2. US Army Research Lab

ICLR, April 2020

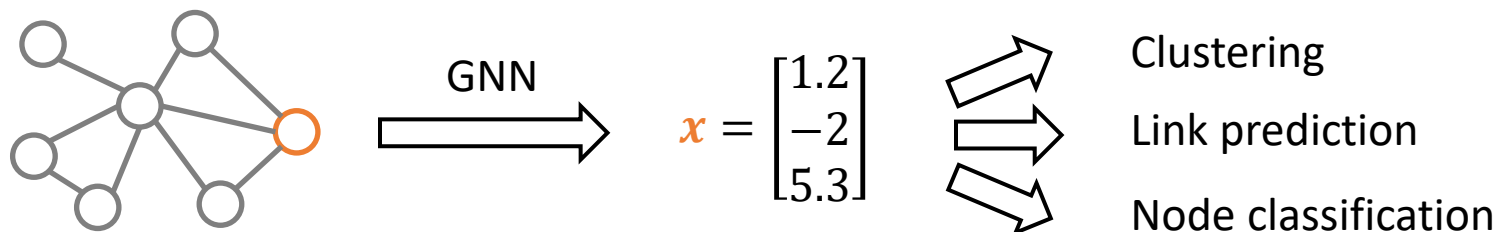
[https://github.com/GraphSAINT/
fpga.usc.edu](https://github.com/GraphSAINT/fpga.usc.edu)



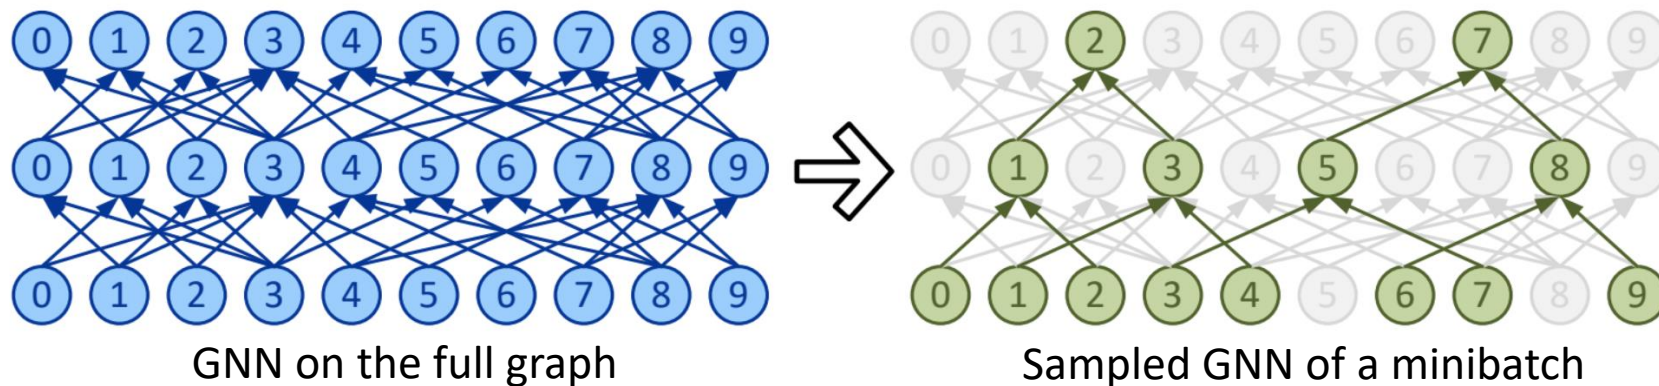
Background: Graph Representation Learning



What it does: Embed each graph node into a low-dimensional vector



Why it is challenging: Num. of multi-hop neighbors explode



Background: Graph Representation Learning

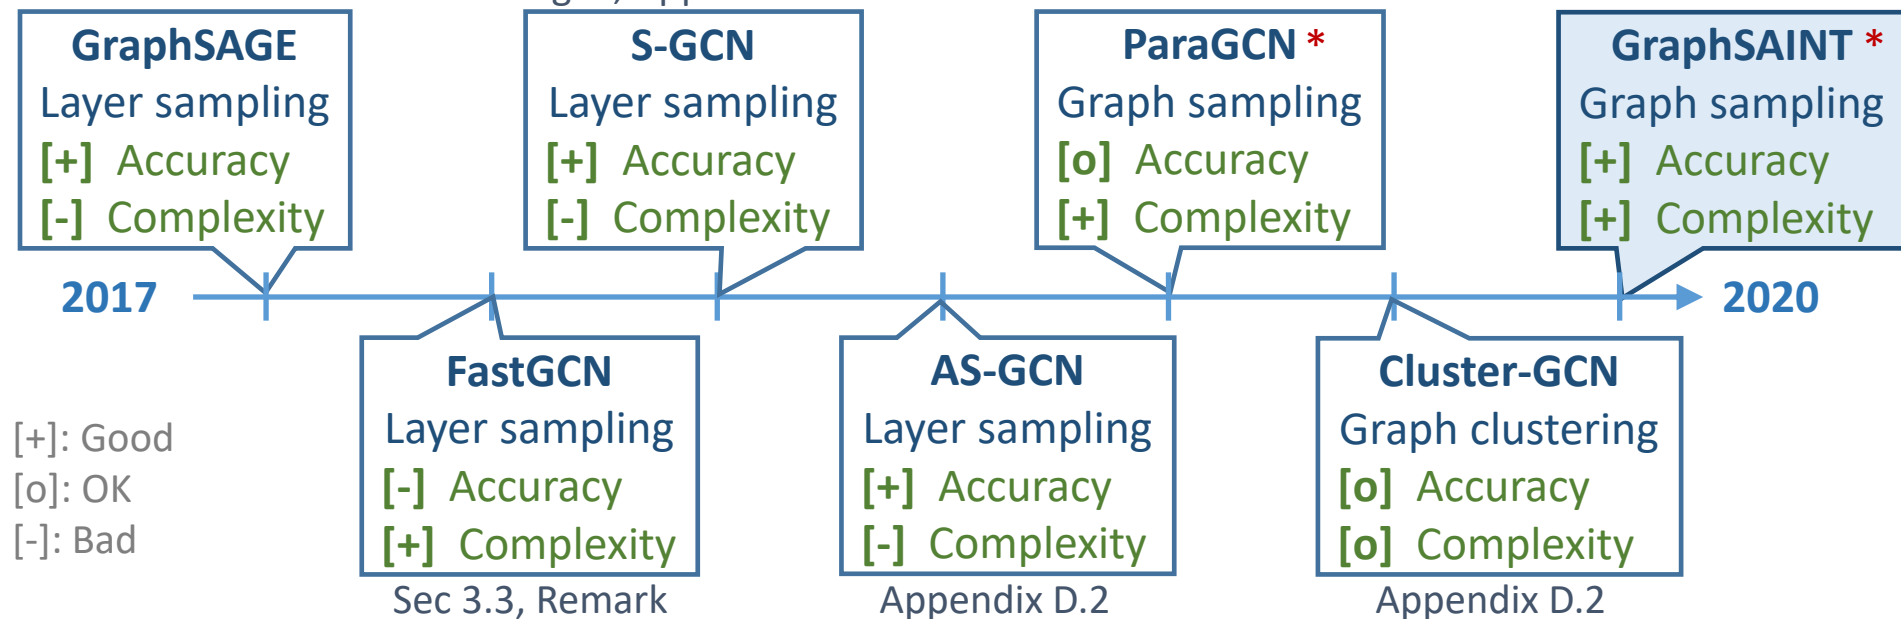


Low training efficiency is now a serious problem in GNN deployment!!

- PinSAGE 2-layer GCN + 7.5billion nodes = **3 days on 16-GPU cluster**
- Deep GNNs Training cost grows **exponentially** with depth

Overview of state-of-the-art minibatch training methods (ref. on slide 11)

Fig. 6, Appendix D.1

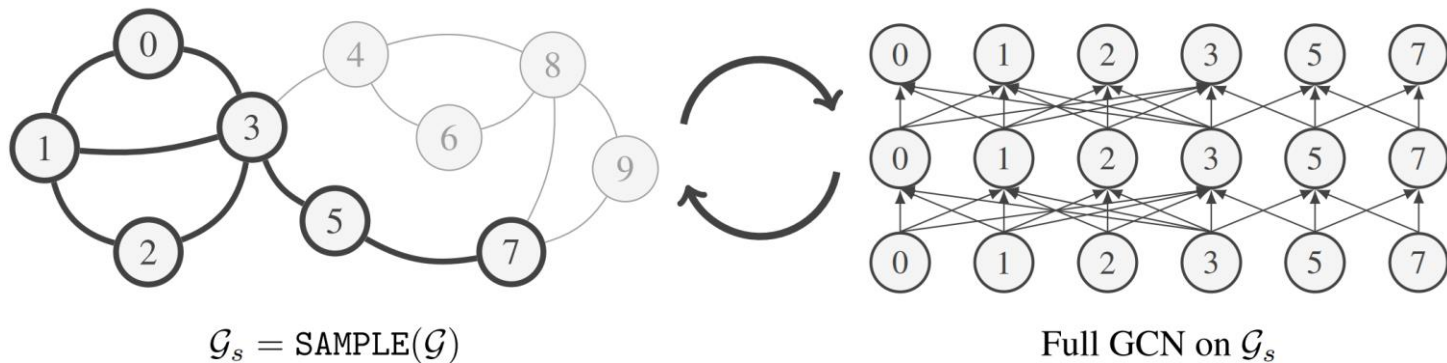




Method: Minibatch Construction

Sample a small subgraph, then build a *complete* GNN

- Constant neighborhood size
- Not an I.I.D. data sampler \rightarrow *Bias elimination*
- Need to preserve connectivity \rightarrow *Variance reduction*



General training framework

- **Graph samplers:** for social networks, molecular graph, traffic networks, ...
- **GNN architectures:** convolution, pooling, attention, skip-connection, ...



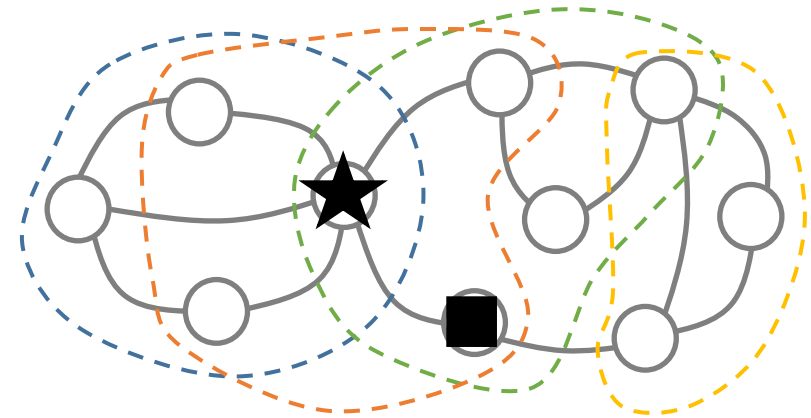
Method: Bias Elimination

Why?

- Popular users will be/need to be sampled more frequently
- Need accurate embedding on everyone

How?

- **Normalize minibatch loss** by probability of sampling each node
- **Normalize neighbor aggregation** by probability of sampling each edge
- **Lightweight preprocessing** to estimate the probability



$$p_{\star} = 3/4$$

$$p_{\star, \blacksquare} = 2/4$$

$$\begin{aligned} Loss_{\text{mini}} &= \frac{1}{|V|} \left(\frac{1}{p_{\star}} Loss_{\star} + \dots \right) \\ x_{\star}' &= \frac{1}{d_{\star}} \left(\frac{1}{p_{\star, \blacksquare}} x_{\blacksquare} + \dots \right) \end{aligned}$$

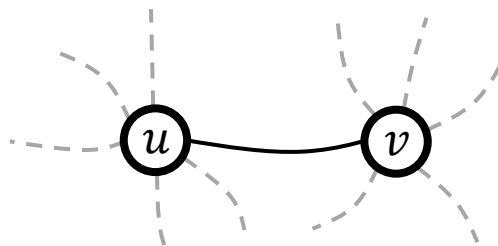


Method: Variance Reduction

How?

- “Important” neighbors to be sampled more frequently

Theorem (independent edge sampling): optimal edge probability for variance minimization is $p_{u,v} \propto \frac{1}{d_u} + \frac{1}{d_v}$.



Sample “**—**” more
Sample “**—**” less

Extensions

- Multi-layer version of random edge sampling \rightarrow Random walk sampler
- Variations of random walk sampler \rightarrow Multi-dimensional RW, etc.



Experiments

Setup

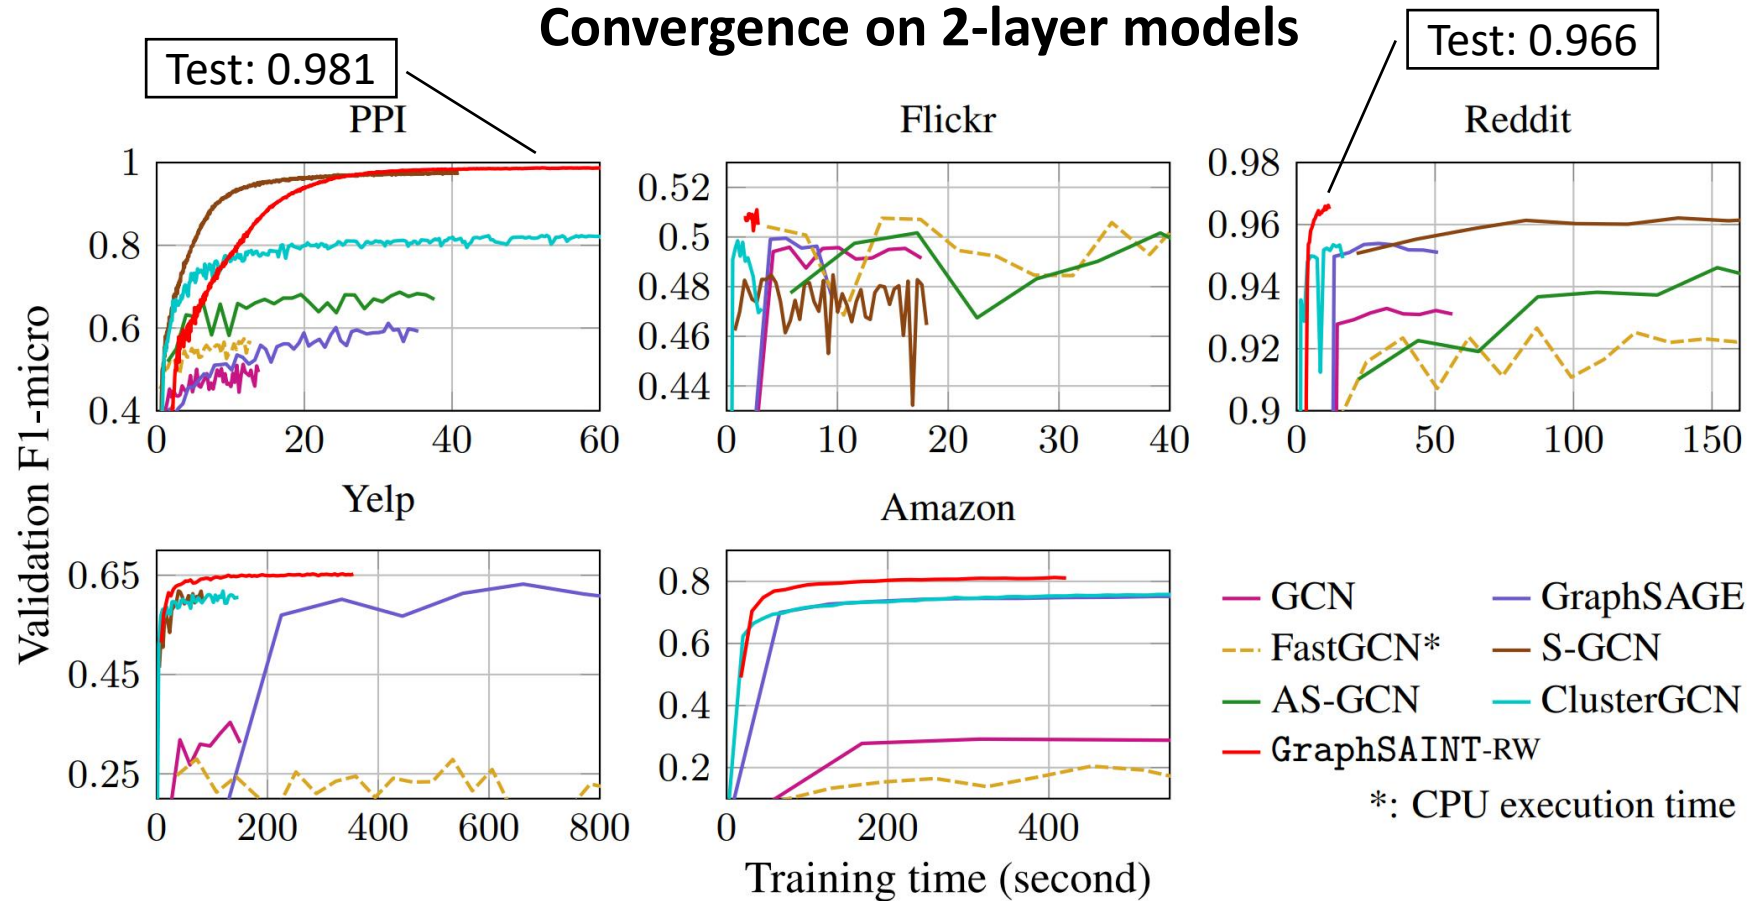
- 5 large graphs Up to 1.5M nodes
- 5 GNN architectures GCN/GraphSAGE, GAT, JK-net, MixHop, GaAN etc.
- 4 graph samplers Node, Edge, Random walk, MD-Random walk
- Multiple depths From 2-layer to 5-layer GNNs

Dataset	Nodes	Edges	Degree	Feature	Classes	Train / Val / Test
PPI	14,755	225,270	15	50	121 (m)	0.66 / 0.12 / 0.22
Flickr	89,250	899,756	10	500	7 (s)	0.50 / 0.25 / 0.25
Reddit	232,965	11,606,919	50	602	41 (s)	0.66 / 0.10 / 0.24
Yelp	716,847	6,977,410	10	300	100 (m)	0.75 / 0.10 / 0.15
Amazon	1,598,960	132,169,734	83	200	107 (m)	0.85 / 0.05 / 0.10
PPI (large version)	56,944	818,716	14	50	121 (m)	0.79 / 0.11 / 0.10



Experiments

Convergence on 2-layer models



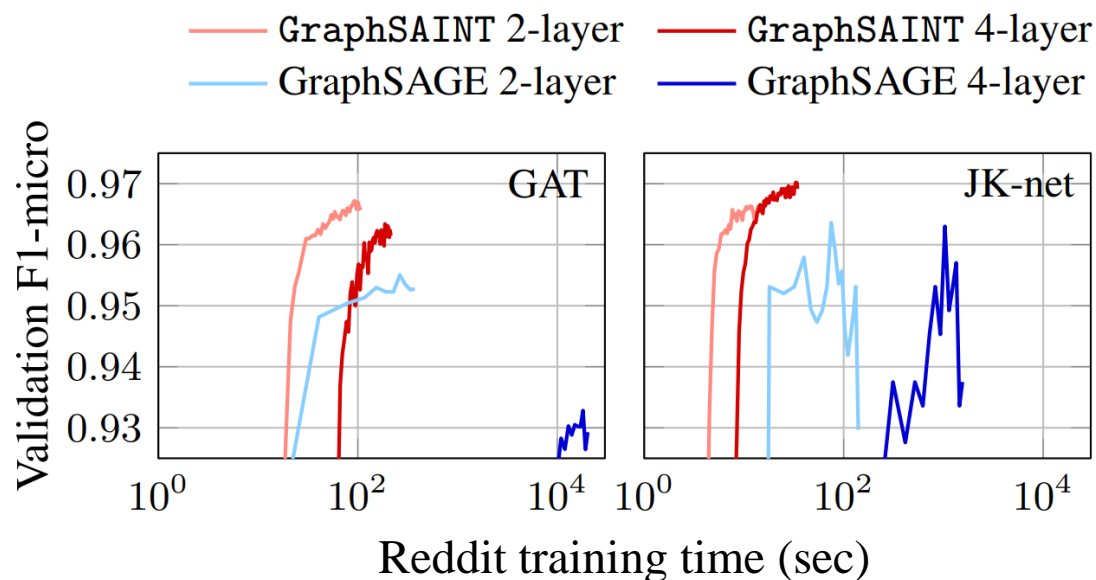
Experiments



Graph clustering
vs.
Graph sampling

	PPI (large version)		Reddit	
	2×512	5×2048	2×128	4×128
ClusterGCN	0.903 ± 0.002	0.994 ± 0.000	0.954 ± 0.001	0.966 ± 0.001
GraphSAINT	0.941 ± 0.003	0.995 ± 0.000	0.966 ± 0.001	0.970 ± 0.001

Layer sampling
vs.
Graph sampling





Conclusion

Minibatch training method for *deep* GNNs on *large* graphs

State-of-the-art training quality (accuracy, speed)

Highly flexible and extensible (graph samplers, GNN architectures)

Official code: <https://github.com/GraphSAINT/GraphSAINT>

 [GraphSAINT](#) / [GraphSAINT](#)

 Watch

5

 Star

63

 Fork

13

PYTORCH
 TensorFlow



Reference implementation: **PyTorch Geometric** library

 [rusty1s](#) / [pytorch_geometric](#)

 Watch

197

 Star

7.3k

 Fork

1.2k

PYTORCH



References

Papers on slide 2

- **GraphSAGE**: “Inductive representation learning on large graphs”. In NeurIPS 2017.
- **FastGCN**: “FastGCN: Fast learning with graph convolutional networks via importance sampling”. In ICLR 2018.
- **S-GCN**: “Stochastic training of graph convolutional networks with variance reduction”. In ICML 2018.
- **AS-GCN**: “Adaptive sampling towards fast graph representation learning”. In NeurIPS 2019.
- **ParaGCN**: “Accurate, efficient and scalable graph embedding”. In IEEE/IPDPS 2019.
- **Cluster-GCN**: “Cluster-GCN: An efficient algorithm for training deep and large graph convolutional networks”. In KDD 2019.
- **GraphSAINT**: “GraphSAINT: Graph sampling based inductive learning method”. In ICLR 2020.



Thank you!

